

Vector Calculus II

Jo Ciucă

Australian National University

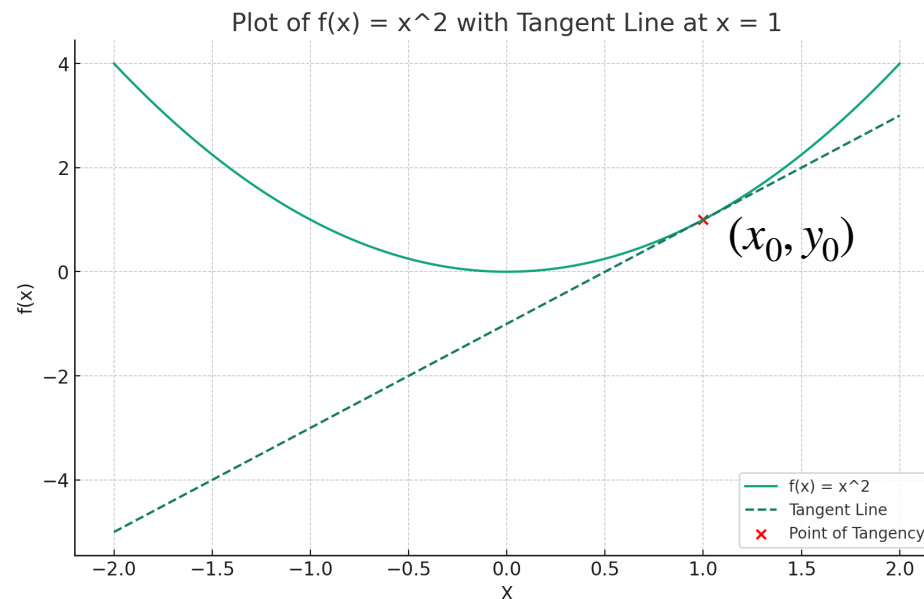
comp36706670@anu.edu.au

Outline

- The Gradient and the Chain Rule
- The Hessian
- Gradients of **Vector-Valued** Functions
- The Jacobian
- Exercises (inc. Gradients of Matrices, iPad session)
- Useful identities for computing gradients

The Derivative: geometrical perspective

- Consider univariate functions $f(x) = y$
- Derivative at a point x_0 is the **slope** of the tangent line at x_0
- Negative slope: f is decreasing
- Positive slope: f is increasing
- The **steeper** the slope (i.e. the larger the absolute value of the slope), the **larger the rate of change of f** .




The Gradient of multivariate functions

- $f: \mathbb{R}^n \rightarrow \mathbb{R}, \mathbf{x} \mapsto f(\mathbf{x}), \mathbf{x} \in \mathbb{R}^n$
- We find the gradient of the function f with respect to \mathbf{x} by
 - varying one variable at a time and keeping the others constant.
 - The gradient is the collection of the partial derivatives.
- We collect the partial derivatives in the row vector


$$\nabla_{\mathbf{x}} f = \text{grad } f = \frac{df}{d\mathbf{x}} = \begin{bmatrix} \frac{\partial f(\mathbf{x})}{\partial x_1} & \frac{\partial f(\mathbf{x})}{\partial x_2} & \dots & \frac{\partial f(\mathbf{x})}{\partial x_n} \end{bmatrix} \in \mathbb{R}^{1 \times n}$$

The Gradient: geometrical perspective

$$f(x, y) = e^{-(x^2+y^2)}$$



Gradient perpendicular to level curve (isocurve).



The Gradient: why we like it

- Encodes how our function responds to changes in the input at a specific point.
- In other words, how “sensitive” our function is to changes in input.
- **Direction:** The gradient points in the direction of the steepest ascent of the function.
- **Magnitude:** The length of the gradient represents the rate of change of the function at that point.
- It likes to be dotted.

$$\left[\frac{\partial f}{\partial x_1} \quad \dots \quad \frac{\partial f}{\partial x_n} \right] \left[\begin{array}{c} \frac{dx_1(t)}{dt} \\ \frac{dx_2(t)}{dt} \\ \vdots \\ \frac{dx_n(t)}{dt} \end{array} \right]$$

Chain Rule



Chain Rule

- Consider a function $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ of two variables x_1 and x_2 .
- $x_1(t)$ and $x_2(t)$ are themselves functions of t .

Approximation $\Delta f \approx f_{x_1} \Delta x_1 + f_{x_2} \Delta x_2$

$$\frac{\Delta f}{\Delta t} \approx f_{x_1} \frac{\Delta x_1}{\Delta t} + f_{x_2} \frac{\Delta x_2}{\Delta t}$$

when Δt goes to 0 : $\frac{df}{dt} = f_{x_1} \frac{dx_1}{dt} + f_{x_2} \frac{dx_2}{dt}$

- Using the chain rule:

$$\frac{df}{dt} = \frac{\partial f}{\partial \mathbf{x}} \frac{d\mathbf{x}}{dt} = \begin{bmatrix} \frac{\partial f}{\partial x_1} & \frac{\partial f}{\partial x_2} \end{bmatrix} \begin{bmatrix} \frac{dx_1(t)}{dt} \\ \frac{dx_2(t)}{dt} \end{bmatrix} = \frac{\partial f}{\partial x_1} \frac{dx_1}{dt} + \frac{\partial f}{\partial x_2} \frac{dx_2}{dt}$$

Chain Rule

- If $f(x_1, x_2)$ is a function of x_1 and x_2 , where $f: \mathbb{R}^2 \rightarrow \mathbb{R}$, $x_1(s, t)$ and $x_2(s, t)$ are themselves functions of two variables s and t , the chain rule yields the **partial derivatives**:

$$\frac{\partial f}{\partial s} = \frac{\partial f}{\partial x_1} \frac{\partial x_1}{\partial s} + \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial s}$$

$$\frac{\partial f}{\partial t} = \frac{\partial f}{\partial x_1} \frac{\partial x_1}{\partial t} + \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial t}$$

- The **gradient** can be obtained by matrix multiplication.

$$\frac{df}{d(s, t)} = \frac{\partial f}{\partial \mathbf{x}} \frac{\partial \mathbf{x}}{\partial (s, t)} = \underbrace{\begin{bmatrix} \frac{\partial f}{\partial x_1} & \frac{\partial f}{\partial x_2} \end{bmatrix}}_{\frac{\partial f}{\partial \mathbf{x}}} \underbrace{\begin{bmatrix} \frac{\partial x_1}{\partial s} & \frac{\partial x_1}{\partial t} \\ \frac{\partial x_2}{\partial s} & \frac{\partial x_2}{\partial t} \end{bmatrix}}_{\frac{\partial \mathbf{x}}{\partial (s, t)}} = \begin{bmatrix} \frac{\partial f}{\partial s} & \frac{\partial f}{\partial t} \end{bmatrix}$$

The gradient likes to be dotted. $= \frac{\partial f}{\partial \mathbf{x}}$

The Hessian Matrix

- Consider a function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ of two variables x_1 and x_2 .
- We consider the second-order partial derivatives, for which:

$$\frac{\partial^2 f}{\partial x_1 \partial x_2} = \frac{\partial^2 f}{\partial x_2 \partial x_1}$$

- The Hessian matrix is the collection of these second-order partial derivatives.

$$H(f) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} \end{bmatrix}$$

- The Hessian measures the **local curvature** at some point (x, y) .
- The gradient tells us about the local slope, i.e. steepness of function.
- The Hessian tells us how the slope is changing, so in a sense is the “derivative of the slope.”

The Hessian Matrix

- For $f: \mathbb{R}^n \rightarrow \mathbb{R}, \mathbf{x} \mapsto f(\mathbf{x}), \mathbf{x} \in \mathbb{R}^n$.

$$\mathbf{H}(f) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

The Hessian is symmetric.

Remember for a square matrix \mathbf{A} :

PD: $x^T \mathbf{A} x > 0$

PSD: $x^T \mathbf{A} x \geq 0$

ND: $x^T \mathbf{A} x < 0$

NSD: $x^T \mathbf{A} x \leq 0$

- The interplay between the Hessian matrix and its **definiteness** properties is profound.
- If Hessian is positive definite (PD) at a point, the function is locally convex. If critical point, then the point is local minimum.
- If negative definite (ND), then the function is locally concave. If critical point, then is local maximum.

Gradients of Vector-Valued Functions

- We discussed partial derivatives and gradients of function $f: \mathbb{R}^n \rightarrow \mathbb{R}$
- We will generalize the concept of the gradient to **vector-valued functions** (vector fields) $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$, where $n \geq 1$ and $m > 1$.
- For a function $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$ and a vector $\mathbf{x} = [x_1, \dots, x_n]^T \in \mathbb{R}^n$, the corresponding vector of function values is given as

$$f(\mathbf{x}) = \begin{bmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{bmatrix} \in \mathbb{R}^m$$

- Writing the vector-valued function in this way allows us to view a vector-valued function $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$ as a vector of functions $[f_1, \dots, f_m]^T$, $f_i: \mathbb{R}^n \rightarrow \mathbb{R}$ that map onto \mathbb{R} .
- The differentiation rules for every f_i are exactly the ones we discussed before.

Gradients of Vector-Valued Functions

- The partial derivative of a vector-valued function $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$ with respect to $x_i \in \mathbb{R}, i = 1, \dots, n$, is given as the vector

$$\frac{\partial f}{\partial x_i} = \begin{bmatrix} \frac{\partial f_1}{\partial x_i} \\ \vdots \\ \frac{\partial f_m}{\partial x_i} \end{bmatrix} = \begin{bmatrix} \lim_{h \rightarrow 0} \frac{f_1(x_1, \dots, x_{i-1}, x_i + h, x_{i+1}, \dots, x_n) - f_1(x)}{h} \\ \vdots \\ \lim_{h \rightarrow 0} \frac{f_m(x_1, \dots, x_{i-1}, x_i + h, x_{i+1}, \dots, x_n) - f_m(x)}{h} \end{bmatrix} \in \mathbb{R}^m$$

- In above, every partial derivative $\frac{\partial f}{\partial x_i}$ is a column vector.
- To obtain the gradient of $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$ with respect to $x \in \mathbb{R}^n$ we collect these partial derivatives:

$$\frac{df(x)}{dx} = \left[\frac{\partial f(x)}{\partial x_1} \quad \dots \quad \frac{\partial f(x)}{\partial x_n} \right] = \begin{bmatrix} \frac{\partial f_1(x)}{\partial x_1} & \dots & \frac{\partial f_1(x)}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m(x)}{\partial x_1} & \dots & \frac{\partial f_m(x)}{\partial x_n} \end{bmatrix} \in \mathbb{R}^{m \times n}$$

$$\mathbf{f}(\mathbf{x}) = \begin{bmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{bmatrix} \in \mathbb{R}^m$$

$$\mathbf{f}(\mathbf{x}) = \begin{bmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{bmatrix} \in \mathbb{R}^m$$

$$\frac{\partial f}{\partial x_i} = \begin{bmatrix} \frac{\partial f_1}{\partial x_i} \\ \vdots \\ \frac{\partial f_m}{\partial x_i} \end{bmatrix}$$

$$\nabla_{\mathbf{x}} f_i = \left[\frac{\partial f_i}{\partial x_1} \quad \frac{\partial f_i}{\partial x_2} \quad \cdots \quad \frac{\partial f_i}{\partial x_n} \right] \in \mathbb{R}^{1 \times n}$$

$$\frac{d\mathbf{f}(\mathbf{x})}{d\mathbf{x}} = \left[\frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_1} \quad \cdots \quad \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_n} \right] =$$

$$\begin{bmatrix} \frac{\partial f_1(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_1(\mathbf{x})}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_m(\mathbf{x})}{\partial x_n} \end{bmatrix} = \begin{bmatrix} \nabla f_1 \\ \nabla f_2 \\ \vdots \\ \nabla f_m \end{bmatrix}$$

The Jacobian

The Jacobian

- The collection of all first-order partial derivatives of a **vector-valued** function $\mathbf{f}: \mathbb{R}^n \rightarrow \mathbb{R}^m$ is called the **Jacobian**. The Jacobian \mathbf{J} is an $m \times n$ matrix, which we define and arrange as follows:

$$\mathbf{J} = \nabla_{\mathbf{x}} \mathbf{f} = \frac{d\mathbf{f}(\mathbf{x})}{d\mathbf{x}} = \begin{bmatrix} \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_n} \end{bmatrix}$$


$$= \begin{bmatrix} \frac{\partial f_1(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_1(\mathbf{x})}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_m(\mathbf{x})}{\partial x_n} \end{bmatrix}$$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}, \quad \mathbf{J}(i, j) = \frac{\partial f_i}{\partial x_j}$$

- The elements of \mathbf{f} define the rows and the elements of \mathbf{x} define the columns of the corresponding Jacobian
- Special case: for a function $\mathbf{f}: \mathbb{R}^n \rightarrow \mathbb{R}^1$ which maps a vector $\mathbf{x} \in \mathbb{R}^n$ onto a scalar, i.e., $m = 1$, the Jacobian is a row vector of dimension $1 \times n$.

To note

- If $f: \mathbb{R} \rightarrow \mathbb{R}$, the gradient is a scalar
- If $f: \mathbb{R}^D \rightarrow \mathbb{R}$, the gradient is a $1 \times D$ row vector
- If $f: \mathbb{R} \rightarrow \mathbb{R}^E$, the gradient is a $E \times 1$ column vector
- If $f: \mathbb{R}^D \rightarrow \mathbb{R}^E$, the gradient is an $E \times D$ matrix



Example - Gradient of a Vector-Valued Function

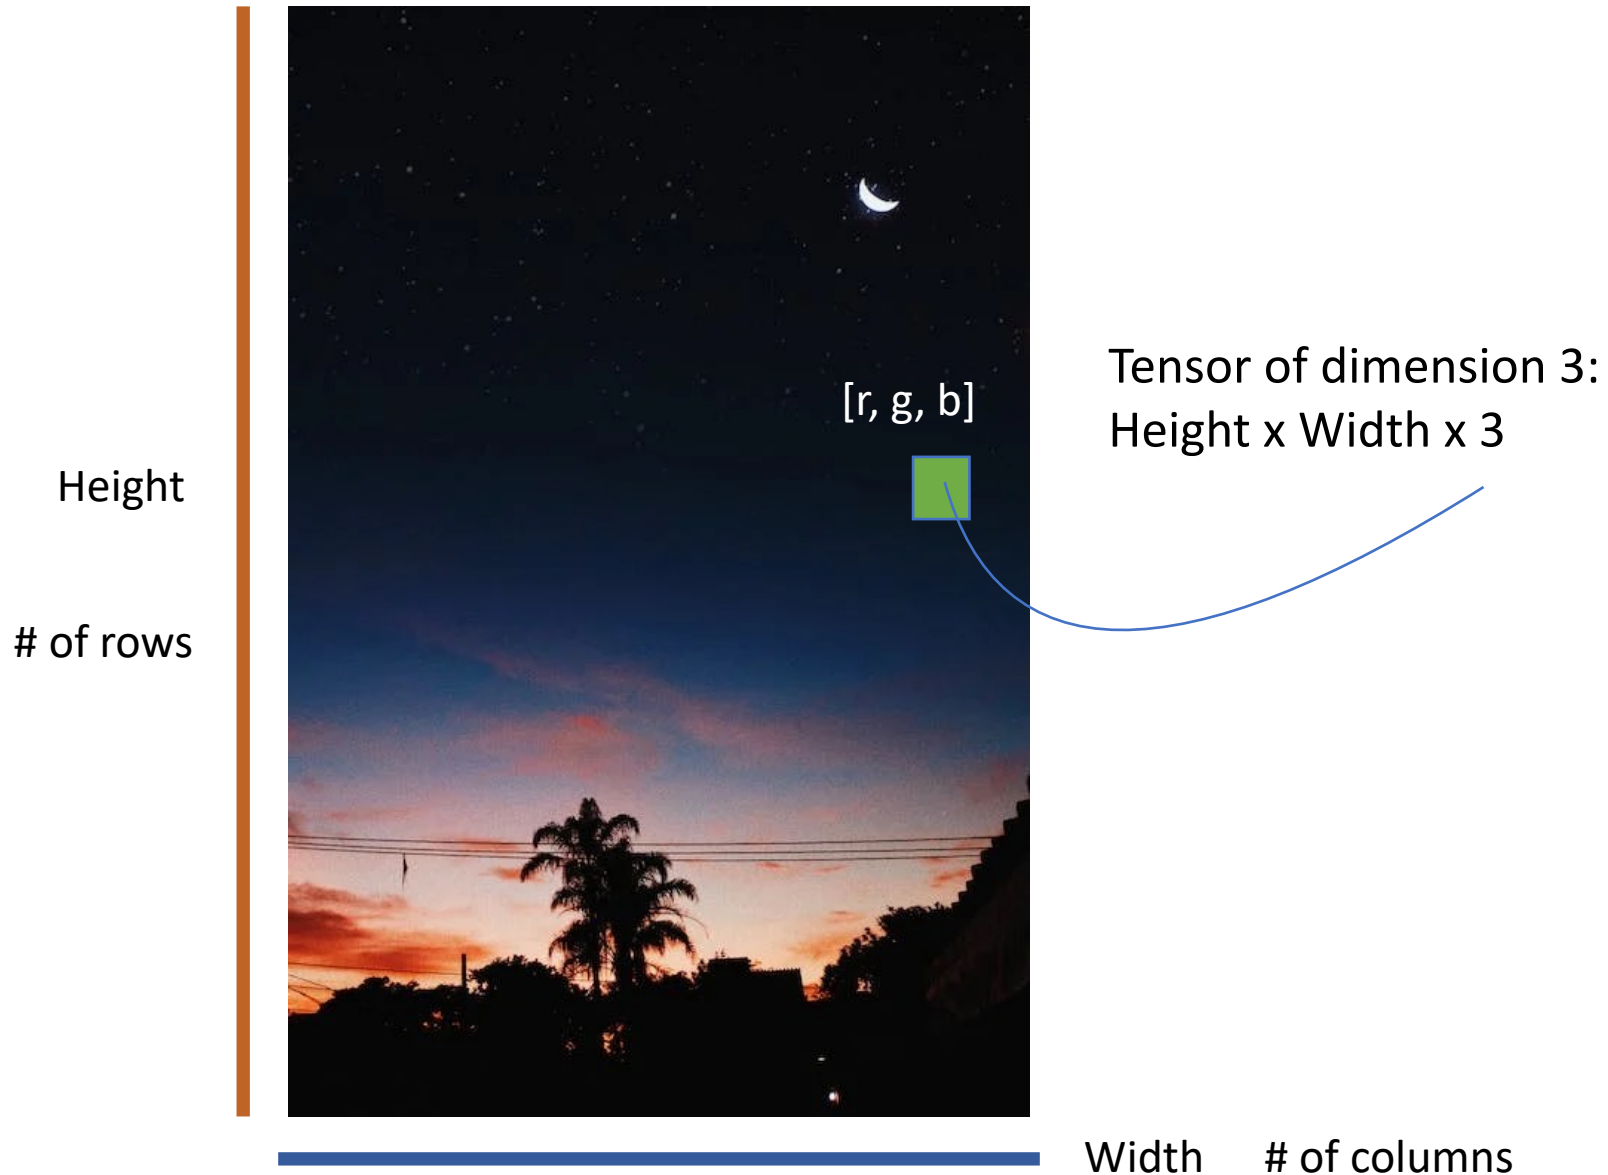
- We are given $\mathbf{f}(\mathbf{x}) = \mathbf{A}\mathbf{x}$, $\mathbf{f}(\mathbf{x}) \in \mathbb{R}^M$, $\mathbf{A} \in \mathbb{R}^{M \times N}$, $\mathbf{x} \in \mathbb{R}^N$.
- To compute the gradient $d\mathbf{f}/d\mathbf{x}$ we first determine the dimension of $d\mathbf{f}/d\mathbf{x}$: Since $\mathbf{f}: \mathbb{R}^N \rightarrow \mathbb{R}^M$, it follows that $d\mathbf{f}/d\mathbf{x} \in \mathbb{R}^{M \times N}$.
- Then, we determine the partial derivatives of \mathbf{f} with respect to every x_j :

$$f_i(\mathbf{x}) = \sum_{j=1}^N A_{ij}x_j \Rightarrow \frac{\partial f_i}{\partial x_j} = A_{ij}$$

- We collect the partial derivatives in the Jacobian and obtain the gradient

$$\frac{\partial \mathbf{f}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_N} \\ \vdots & & \vdots \\ \frac{\partial f_M}{\partial x_1} & \cdots & \frac{\partial f_M}{\partial x_N} \end{bmatrix} = \begin{bmatrix} A_{11} & \cdots & A_{1N} \\ \vdots & & \vdots \\ A_{M1} & \cdots & A_{MN} \end{bmatrix} = \mathbf{A} \in \mathbb{R}^{M \times N}$$

What even is a tensor? The machine learning answer.



Example #1 - Chain Rule

- Consider the function $h: \mathbb{R} \rightarrow \mathbb{R}$, $h(t) = (f \circ g)(t)$ with

$$f: \mathbb{R}^2 \rightarrow \mathbb{R}$$

$$g: \mathbb{R} \rightarrow \mathbb{R}^2$$

$$f(\mathbf{x}) = \exp(x_1 x_2^2)$$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = g(t) = \begin{bmatrix} t \cos t \\ t \sin t \end{bmatrix}$$

- We compute the gradient of h with respect to t . Since $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ and $g: \mathbb{R} \rightarrow \mathbb{R}^2$ we note that

$$\frac{\partial f}{\partial \mathbf{x}} \in \mathbb{R}^{1 \times 2}, \quad \frac{\partial g}{\partial t} \in \mathbb{R}^{2 \times 1}$$

- The desired gradient is computed by applying the chain rule:

$$\begin{aligned} \frac{dh}{dt} &= \frac{\partial f}{\partial \mathbf{x}} \frac{\partial \mathbf{x}}{\partial t} = \begin{bmatrix} \frac{\partial f}{\partial x_1} & \frac{\partial f}{\partial x_2} \end{bmatrix} \begin{bmatrix} \frac{\partial x_1}{\partial t} \\ \frac{\partial x_2}{\partial t} \end{bmatrix} \\ &= \begin{bmatrix} \exp(x_1 x_2^2) x_2^2 & 2 \exp(x_1 x_2^2) x_1 x_2 \end{bmatrix} \begin{bmatrix} \cos t - t \sin t \\ \sin t + t \cos t \end{bmatrix} \\ &= \exp(x_1 x_2^2) (x_2^2 (\cos t - t \sin t) + 2 x_1 x_2 (\sin t + t \cos t)) \\ &\text{where } x_1 = t \cos t \text{ and } x_2 = t \sin t \end{aligned}$$

Example #2 - Gradient of a Least-Squares Loss in a Linear Model

- Let us consider the linear model

$$\mathbf{y} = \Phi \theta$$

where $\theta \in \mathbb{R}^D$ is a parameter vector, $\Phi \in \mathbb{R}^{N \times D}$ are input features and $\mathbf{y} \in \mathbb{R}^N$ are the corresponding observations. We define the functions

$$L(e) := \|e\|^2,$$

$$e(\theta) := \mathbf{y} - \Phi \theta$$

- We seek $\frac{\partial L}{\partial \theta}$, and we will use the chain rule for this purpose. L is called a least-squares loss function.
- First, we determine the dimensionality of the gradient as

$$\frac{\partial L}{\partial \theta} \in \mathbb{R}^{1 \times D}$$

- The chain rule allows us to compute the gradient as

$$\frac{\partial L}{\partial \theta} = \frac{\partial L}{\partial e} \frac{\partial e}{\partial \theta}$$

Example #2 - Gradient of a Least-Squares Loss in a Linear Model

- We know that $\|e\|^2 = e^T e$ and determine

$$\frac{\partial L}{\partial e} = 2e^T \in \mathbb{R}^{1 \times N}$$

- Further, we obtain

$$\frac{\partial e}{\partial \theta} = -\Phi \in \mathbb{R}^{N \times D}$$

- Our desired derivative is

$$\frac{\partial L}{\partial \theta} = -2e^T \Phi = - \underbrace{2(y^T - \theta^T \Phi^T)}_{1 \times N} \underbrace{\Phi}_{N \times D} \in \mathbb{R}^{1 \times D}$$

Example #3: Gradients of Matrices

- Consider the following:

$$\mathbf{f} = \mathbf{A}\mathbf{x}, \quad \mathbf{f} \in \mathbb{R}^M, \quad \mathbf{A} \in \mathbb{R}^{M \times N}, \quad \mathbf{x} \in \mathbb{R}^N$$

- We seek the gradient $\frac{d\mathbf{f}}{d\mathbf{A}}$

- First, we determine the dimension of the gradient

$$\frac{d\mathbf{f}}{d\mathbf{A}} \in \mathbb{R}^{M \times (M \times N)}$$

- By definition, the gradient is the collection of the partial derivatives:

$$\frac{d\mathbf{f}}{d\mathbf{A}} = \begin{bmatrix} \frac{\partial f_1}{\partial \mathbf{A}} \\ \vdots \\ \frac{\partial f_M}{\partial \mathbf{A}} \end{bmatrix}, \quad \frac{\partial f_i}{\partial \mathbf{A}} \in \mathbb{R}^{1 \times (M \times N)}$$

- To compute the partial derivatives, we explicitly write out the matrix vector multiplication

$$f_i = \sum_{j=1}^N A_{ij}x_j, \quad i = 1, \dots, M,$$

$$f_i = \sum_{j=1}^N A_{ij} x_j, \quad i = 1, \dots, M$$

- The partial derivatives are then given as

$$\frac{\partial f_i}{\partial A_{iq}} = x_q$$

- Partial derivatives of f_i with respect to a row of \mathbf{A} are given as

$$\frac{\partial f_i}{\partial A_{i,:}} = \mathbf{x}^T \in \mathbb{R}^{1 \times 1 \times N}, \quad \frac{\partial f_i}{\partial A_{k \neq i,:}} = \mathbf{0}^T \in \mathbb{R}^{1 \times 1 \times N}$$

- Since f_i maps onto \mathbb{R} and each row of \mathbf{A} is of size $1 \times N$, we obtain a $1 \times 1 \times N$ sized tensor as the partial derivative of f_i with respect to a row of \mathbf{A} .
- We stack the partial derivatives and get the desired gradient

$$\frac{\partial f_i}{\partial \mathbf{A}} = \begin{bmatrix} \mathbf{0}^T \\ \vdots \\ \mathbf{0}^T \\ \mathbf{x}^T \\ \mathbf{0}^T \\ \vdots \\ \mathbf{0}^T \end{bmatrix} \in \mathbb{R}^{1 \times (M \times N)}$$

Example #4: Gradient of Matrices with Respect to Matrices

- Consider a matrix $\mathbf{R} \in \mathbb{R}^{M \times N}$ and $f: \mathbb{R}^{M \times N} \rightarrow \mathbb{R}^{N \times N}$ with

$$f(\mathbf{R}) = \mathbf{R}^T \mathbf{R} =: \mathbf{K} \in \mathbb{R}^{N \times N}$$

- We seek the gradient $\frac{d\mathbf{K}}{d\mathbf{R}}$

- First, the dimension of the gradient is given as

$$\frac{d\mathbf{K}}{d\mathbf{R}} \in \mathbb{R}^{(N \times N) \times (M \times N)}$$

$$\frac{dK_{pq}}{d\mathbf{R}} \in \mathbb{R}^{1 \times M \times N}$$

for $p, q = 1, \dots, N$, where K_{pq} is the pq th entry of $\mathbf{K} = f(\mathbf{R})$.

- Denoting the i th column of \mathbf{R} by \mathbf{r}_i , every entry of \mathbf{K} is given by the dot product of two columns of \mathbf{R} , i.e.,

$$K_{pq} = \mathbf{r}_p^T \mathbf{r}_q = \sum_{m=1}^M R_{mp} R_{mq}$$

Example #4: Gradient of Matrices with Respect to Matrices

- Denoting the i th column of \mathbf{R} by \mathbf{r}_i , every entry of \mathbf{K} is given by the dot product of two columns of \mathbf{R} , i.e.,

$$K_{pq} = \mathbf{r}_p^T \mathbf{r}_q = \sum_{m=1}^M R_{mp} R_{mq}$$

- We now compute the partial derivative $\frac{\partial K_{pq}}{\partial R_{ij}}$, we obtain

$$\frac{\partial K_{pq}}{\partial R_{ij}} = \sum_{m=1}^M \frac{\partial}{\partial R_{ij}} R_{mp} R_{mq} = \partial_{pqij}$$

$$\partial_{pqij} = \begin{cases} R_{iq} & \text{if } j = p, p \neq q \\ R_{ip} & \text{if } j = q, p \neq q \\ 2R_{iq} & \text{if } j = p, p = q \\ 0 & \text{otherwise} \end{cases}$$

- The desired gradient has the dimension $(N \times N) \times (M \times N)$, and every single entry of this tensor is given by ∂_{pqij} , where $p, q, j = 1, \dots, N$ and $i = 1, \dots, M$

Useful Identities for Computing Gradients

- Some useful gradients that are frequently required in machine learning.

Note that the trace of a square matrix, $\text{tr}(A) = \sum_{i=1}^n A_{ii}$.

$$\frac{\partial \mathbf{x}^T \mathbf{a}}{\partial \mathbf{x}} = \mathbf{a}^T$$

$$\frac{\partial \mathbf{a}^T \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}^T$$

You should be able to calculate these gradients.

$$\frac{\partial \mathbf{a}^T \mathbf{X} \mathbf{b}}{\partial \mathbf{X}} = \mathbf{a} \mathbf{b}^T$$

$$\frac{\partial \mathbf{x}^T \mathbf{B} \mathbf{x}}{\partial \mathbf{x}} = \mathbf{x}^T (\mathbf{B} + \mathbf{B}^T)$$

$$\frac{\partial}{\partial \mathbf{s}} (\mathbf{x} - \mathbf{A} \mathbf{s})^T \mathbf{W} (\mathbf{x} - \mathbf{A} \mathbf{s}) = -2(\mathbf{x} - \mathbf{A} \mathbf{s})^T \mathbf{W} \mathbf{A} \quad \text{for symmetric } \mathbf{W}$$